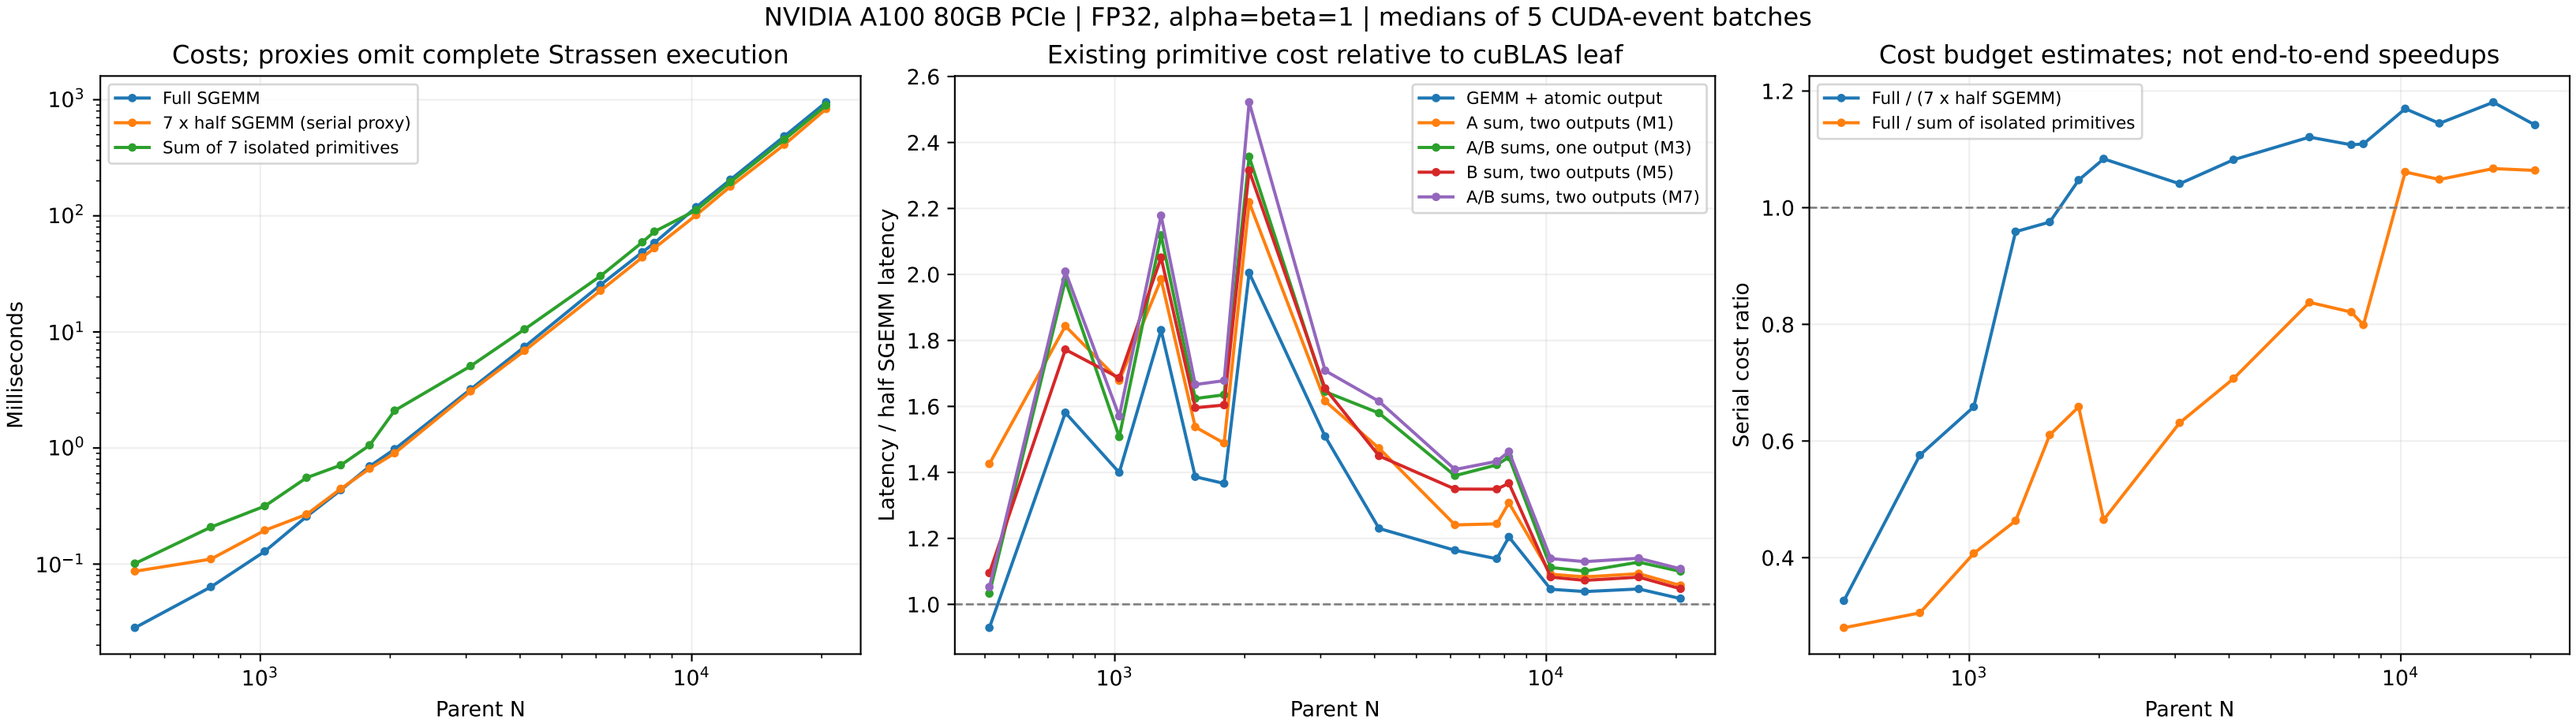
NVIDIA A100 80GB PCIe | FP32, alpha=beta=1 | medians of 5 CUDA-event batches
Costs; proxies omit complete Strassen execution
Milliseconds
Parent N
Full SGEMM
7 x half SGEMM (serial proxy)
Sum of 7 isolated primitives
Existing primitive cost relative to cuBLAS leaf
Latency / half SGEMM latency
Parent N
GEMM + atomic output
A sum, two outputs (M1)
A/B sums, one output (M3)
B sum, two outputs (M5)
A/B sums, two outputs (M7)
Cost budget estimates; not end-to-end speedups
Serial cost ratio
Parent N
Full / (7 x half SGEMM)
Full / sum of isolated primitives